

Problem Statement

A US-based housing company named Surprise Housing has decided to enter the Australian market. The company uses data analytics to purchase houses at a price below their actual values and flip them on at a higher price. For the same purpose, the company has collected a data set from the sale of houses in Australia. The data is provided in the CSV file below.

The company is looking at prospective properties to buy to enter the market. You are required to build a regression model using regularisation in order to predict the actual value of the prospective properties and decide whether to invest in them or not.

The company wants to know:

- Which variables are significant in predicting the price of a house, and
- How well those variables describe the price of a house.

Also, determine the optimal value of lambda for ridge and lasso regression.

Business Goal

You are required to model the price of houses with the available independent variables. This model will then be used by the management to understand how exactly the prices vary with the variables. They can accordingly manipulate the strategy of the firm and concentrate on areas that will yield high returns. Further, the model will be a good way for management to understand the pricing dynamics of a new market.

```
In [531]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np

from sklearn.preprocessing import StandardScaler

from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import KFold
from sklearn.model_selection import GridSearchCV

from sklearn.linear_model import Ridge
from sklearn.linear_model import Lasso
from sklearn.linear_model import LinearRegression

from sklearn.pipeline import make_pipeline
from sklearn import metrics

# Data display customization
pd.set_option('display.max_rows', 2000)
pd.set_option('display.max_columns', 100)

# hide warnings
import warnings
warnings.filterwarnings('ignore')
```

Data Visualization and Understanding

```
In [532]: housing_df = pd.read_csv("train.csv")  
housing_df.shape
```

```
Out[532]: (1460, 81)
```

```
In [453]: housing_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1460 entries, 0 to 1459
Data columns (total 81 columns):
Id                1460 non-null int64
MSSubClass        1460 non-null int64
MSZoning          1460 non-null object
LotFrontage       1201 non-null float64
LotArea           1460 non-null int64
Street            1460 non-null object
Alley             91 non-null object
LotShape          1460 non-null object
LandContour       1460 non-null object
Utilities         1460 non-null object
LotConfig         1460 non-null object
LandSlope         1460 non-null object
Neighborhood      1460 non-null object
Condition1        1460 non-null object
Condition2        1460 non-null object
BldgType          1460 non-null object
HouseStyle        1460 non-null object
OverallQual       1460 non-null int64
OverallCond       1460 non-null int64
YearBuilt         1460 non-null int64
YearRemodAdd      1460 non-null int64
RoofStyle         1460 non-null object
RoofMatl          1460 non-null object
Exterior1st       1460 non-null object
Exterior2nd       1460 non-null object
MasVnrType        1452 non-null object
MasVnrArea        1452 non-null float64
ExterQual         1460 non-null object
ExterCond         1460 non-null object
Foundation        1460 non-null object
BsmtQual          1423 non-null object
BsmtCond          1423 non-null object
BsmtExposure      1422 non-null object
BsmtFinType1      1423 non-null object
BsmtFinSF1        1460 non-null int64
BsmtFinType2      1422 non-null object
BsmtFinSF2        1460 non-null int64
BsmtUnfSF         1460 non-null int64
TotalBsmtSF       1460 non-null int64
Heating           1460 non-null object
HeatingQC         1460 non-null object
CentralAir        1460 non-null object
Electrical        1459 non-null object
1stFlrSF          1460 non-null int64
2ndFlrSF          1460 non-null int64
LowQualFinSF      1460 non-null int64
GrLivArea         1460 non-null int64
BsmtFullBath      1460 non-null int64
BsmtHalfBath      1460 non-null int64
FullBath          1460 non-null int64
HalfBath          1460 non-null int64
BedroomAbvGr      1460 non-null int64
KitchenAbvGr      1460 non-null int64
KitchenQual       1460 non-null object
TotRmsAbvGrd      1460 non-null int64
Functional        1460 non-null object
Fireplaces        1460 non-null int64
FireplaceQu       770 non-null object
GarageType        1379 non-null object
GarageYrBltd      1379 non-null float64
```

```
In [533]: # Checking how many are categorical and how many numeric variable in the given data-set
# to make sure this falls inline with the data-dictionary
housing_categorical=housing_df.select_dtypes(include=['object'])
housing_numeric=housing_df.select_dtypes(include=np.number)
print("categorical_vars:" + str(housing_categorical.columns))
print("numeric_vars:" + str(housing_numeric.columns))


categorical_vars:Index(['MSZoning', 'Street', 'Alley', 'LotShape', 'LandContour', 'Utilities',
                        'LotConfig', 'LandSlope', 'Neighborhood', 'Condition1', 'Condition2',
                        'BldgType', 'HouseStyle', 'RoofStyle', 'RoofMatl', 'Exterior1st',
                        'Exterior2nd', 'MasVnrType', 'ExterQual', 'ExterCond', 'Foundation',
                        'BsmtQual', 'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'BsmtFinType2',
                        'Heating', 'HeatingQC', 'CentralAir', 'Electrical', 'KitchenQual',
                        'Functional', 'FireplaceQu', 'GarageType', 'GarageFinish', 'GarageQual',
                        'GarageCond', 'PavedDrive', 'PoolQC', 'Fence', 'MiscFeature',
                        'SaleType', 'SaleCondition'],
                        dtype='object')
numeric_vars:Index(['Id', 'MSSubClass', 'LotFrontage', 'LotArea', 'OverallQual',
                    'OverallCond', 'YearBuilt', 'YearRemodAdd', 'MasVnrArea', 'BsmtFinSF1',
                    'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF', '1stFlrSF', '2ndFlrSF',
                    'LowQualFinSF', 'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath',
                    'HalfBath', 'BedroomAbvGr', 'KitchenAbvGr', 'TotRmsAbvGrd',
                    'Fireplaces', 'GarageYrBlt', 'GarageCars', 'GarageArea', 'WoodDeckSF',
                    'OpenPorchSF', 'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea',
                    'MiscVal', 'MoSold', 'YrSold', 'SalePrice'],
                    dtype='object')
```

Check if the data is normally distributed against target variable **SalePrice**

```
In [455]: housing_df['SalePrice'].describe().round(2)
```

```
Out[455]: count      1460.0
mean      180921.2
std       79442.5
min       34900.0
25%      129975.0
50%      163000.0
75%      214000.0
max       755000.0
Name: SalePrice, dtype: float64
```

```
In [456]: # from the dist plot it we can see price column is not normally distributed
, the Gaussian curve is scewed towards left
sns.distplot(housing_df['SalePrice'])
#plt.hist(housing_df['SalePrice'])
plt.show()
```



Doing a normality test on target variable `SalePrice` using Shapiro-Wilk normality test from scipy module

```
In [457]: # perform normality test on target variable sale price
from scipy.stats import shapiro
def normality_test_price():
    stat,p = shapiro(housing_df['SalePrice'])
    alpha=0.05
    print('Statistics=0.3f, p=0.3f' % (stat, p))
    if p > alpha:
        print('SalePrice looks following Gaussian (fail to reject Null Hypothesis H0)')
    else:
        print('SalePrice does not look following Gaussian (reject Null Hypothesis H0)')
normality_test_price()
```


Statistics=0.870, p=0.000

SalePrice does not look following Gaussian (reject Null Hypothesis H0)


```
In [458]: # lets see how house price varies based on number of bedroom above garage
housing_df['BedroomAbvGr'].value_counts()
```

```
Out[458]: 3    804
          2    358
          4    213
          1     50
          5     21
          6       7
          0        6
          8         1
          Name: BedroomAbvGr, dtype: int64
```

```
In [326]: sns.boxplot(x='BedroomAbvGr',y='SalePrice', data=housing_df)
plt.show()
```



```
In [327]: sns.barplot(x='BedroomAbvGr',y='SalePrice', data=housing_df)
plt.show()
```




Looks like there could be some data imbalance in the data set. Ideally house price should increase as number of bedroom increases. Here the trend is beyond 4 bedrooms the house price is not creasing. There could be multiple reason for that. One reason could be locality/neighbourhood as well.

```
In [11]: # MSZoning: Identifies the general zoning classification of the sale
housing_df['MSZoning'].value_counts()
```

```
Out[11]: RL      1151
RM        218
FV         65
RH         16
C (all)    10
Name: MSZoning, dtype: int64
```

```
In [15]: # Visualize how zoning relate to the target variable SalePrice
sns.boxplot(x=housing_df['MSZoning'],y=housing_df['SalePrice'])
plt.show()
```



Most of the price distribution based on zoning category are lying within 400K mark. Very few residential property price is significantly high. We might need to consider those as outliers later.

```
In [19]: # Utilities: Type of utilities available
housing_df['Utilities'].value_counts()
```


```
Out[19]: AllPub      1459
         NoSeWa       1
         Name: Utilities, dtype: int64
```



```
In [27]: # Lets see how contision's are effecting sale price if any to understand if
         # this a important factor to decide upon pricing
         plt.figure(figsize=(8, 6))
         plt.subplot(1,2,1)
         sns.barplot(y='SalePrice',x='Condition1',data=housing_df)
         plt.xticks(rotation=80)

         plt.subplot(1,2,2)
         sns.barplot(y='SalePrice',x='Condition2',data=housing_df)
         plt.xticks(rotation=80)


         plt.tight_layout()
         plt.show()
```




```
In [33]: # MSSubClass: Identifies the type of dwelling involved in the sale
         housing_df['MSSubClass'].value_counts()
```

```
Out[33]: 20      536
         60      299
         50      144
         120     87
         30      69
         160     63
         70      60
         80      58
         90      52
         190     30
         85      20
         75      16
         45      12
         180     10
         40       4
         Name: MSSubClass, dtype: int64
```


```
In [34]: sns.distplot(housing_df['MSSubClass'])  
plt.show()
```




```
In [35]: # Neighborhood: Physical locations within Ames city limits  
plt.figure(figsize=(10, 6))  
sns.barplot(y='SalePrice', x='Neighborhood', data=housing_df)  
plt.xticks(rotation=60)  
plt.show()
```




```
In [36]: # BldgType: Type of dwelling
plt.figure(figsize=(8, 6))
sns.barplot(y='SalePrice', x='BldgType', data=housing_df)
plt.xticks(rotation=60)
plt.show()
```



```
In [37]: # HouseStyle: Style of dwelling
plt.figure(figsize=(8, 6))
sns.barplot(y='SalePrice', x='HouseStyle', data=housing_df)
plt.xticks(rotation=60)
plt.show()
```



```
In [38]: # LotArea: Lot size in square feet
# visualise area-price relationship
sns.regplot(x="LotArea", y="SalePrice", data=housing_df, fit_reg=False)
plt.show()
```



Data Exploration

To perform linear regression, the (numeric) target variable should be linearly related to *at least one another numeric variable*. Let's see whether that's true in this case.

We'll first subset the list of all (independent) numeric variables, and then make a **heatmap**

```
In [459]: # all numeric (float and int) variables in the dataset
housing_numeric = housing_df.select_dtypes(include=['float64', 'int64'])
housing_numeric.head()
```

Out[459]:

	Id	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	YearBuilt	YearRemodAdd	MasVnrArea
0	1	60	65.0	8450	7	5	2003	2003	196.0
1	2	20	80.0	9600	6	8	1976	1976	0.0
2	3	60	68.0	11250	7	5	2001	2002	162.0
3	4	70	60.0	9550	7	5	1915	1970	0.0
4	5	60	84.0	14260	8	5	2000	2000	350.0

```
In [460]: housing_numeric.isnull().sum()
```

```
Out[460]: Id                0
MSSubClass                0
LotFrontage              259
LotArea                  0
OverallQual               0
OverallCond               0
YearBuilt                 0
YearRemodAdd              0
MasVnrArea                8
BsmtFinSF1                0
BsmtFinSF2                0
BsmtUnfSF                 0
TotalBsmtSF               0
1stFlrSF                  0
2ndFlrSF                  0
LowQualFinSF              0
GrLivArea                 0
BsmtFullBath              0
BsmtHalfBath              0
FullBath                  0
HalfBath                  0
BedroomAbvGr              0
KitchenAbvGr              0
TotRmsAbvGrd              0
Fireplaces                0
GarageYrBlt               81
GarageCars                0
GarageArea                0
WoodDeckSF                0
OpenPorchSF               0
EnclosedPorch             0
3SsnPorch                 0
ScreenPorch               0
PoolArea                  0
MiscVal                   0
MoSold                    0
YrSold                    0
SalePrice                 0
dtype: int64
```

We can drop `GarageYrBlt` and `LotFrontage` variables as from the data dictionary these does not sounds to impact much in our model

```
In [534]: # dropping GarageYrBlt and LotFrontage along with ID
housing_numeric = housing_numeric.drop(['LotFrontage', 'Id', 'GarageYrBlt'],
axis=1)
```

```
In [493]: housing_numeric['PoolArea'].value_counts()
```

```
Out[493]: 0      1453
          738      1
          648      1
          576      1
          555      1
          519      1
          512      1
          480      1
          Name: PoolArea, dtype: int64
```

```
In [42]: housing_numeric['ScreenPorch'].value_counts()
```

```
Out[42]: 0      1344
          192      6
          224      5
          120      5
          189      4
          ...
          182      1
          440      1
          178      1
          312      1
          480      1
          Name: ScreenPorch, Length: 76, dtype: int64
```

```
In [43]: housing_numeric['3SsnPorch'].value_counts()
```

```
Out[43]: 0      1436
          168      3
          216      2
          144      2
          180      2
          245      1
          238      1
          290      1
          196      1
          182      1
          407      1
          304      1
          162      1
          153      1
          320      1
          140      1
          130      1
          96      1
          23      1
          508      1
          Name: 3SsnPorch, dtype: int64
```

As per data dictionary definition the variables 3SsnPorch , ScreenPorch , PoolArea should have values in feet. But for most the data points these are having 0 feet which suggests either the corresponding property does not have these facilities or it has missing value. Probably we can drop such columns where majority of the value is '0'. These variables won't be significant enough for generalizing the housing price prediction.

```
In [535]: # replacing the value 0 with NaN to find how much percentage of such 0 value variables are there
for col in housing_numeric.columns:
    #print(col)
    housing_numeric[col].replace(0,np.NaN,inplace=True)
#housing_numeric.isnull().sum()
```


```
In [536]: nullcols=(100*(housing_numeric.isnull().sum()/len(housing_numeric.index)).round(2)).sort_values(ascending=False)
nullcols[nullcols>0]
```

```
Out[536]: PoolArea          100.0
3SsnPorch          98.0
LowQualFinSF       98.0
MiscVal            96.0
BsmtHalfBath       94.0
ScreenPorch        92.0
BsmtFinSF2         89.0
EnclosedPorch      86.0
HalfBath           63.0
MasVnrArea         60.0
BsmtFullBath       59.0
2ndFlrSF           57.0
WoodDeckSF         52.0
Fireplaces         47.0
OpenPorchSF        45.0
BsmtFinSF1         32.0
BsmtUnfSF          8.0
GarageCars         6.0
GarageArea         6.0
TotalBsmtSF        3.0
FullBath           1.0
dtype: float64
```

```
In [357]: # for example check the distribution of PoolArea, ScreenPorch, BsmtFinSF2 from original data set
#plt.figure(figsize=(8, 6))
plt.subplot(1,3,1)
sns.distplot(housing_df['PoolArea'])

plt.subplot(1,3,2)
sns.distplot(housing_df['ScreenPorch'])

plt.subplot(1,3,3)
sns.distplot(housing_df['BsmtFinSF2'])
plt.tight_layout()
plt.show()
```



From the above analysis we can clearly drop columns where encoded NaN for corresponding 0 value is more than 89% in data set

```
In [407]: nullcols=nullcols>88
#
```

```
Out[407]: PoolArea      100.0
3SsnPorch      98.0
LowQualFinSF   98.0
MiscVal        96.0
BsmtHalfBath   94.0
ScreenPorch    92.0
BsmtFinSF2     89.0
dtype: float64
```

```
In [537]: housing_numeric.drop(columns=['PoolArea', '3SsnPorch', 'LowQualFinSF', 'MiscVal',
    'BsmtHalfBath', 'ScreenPorch', 'BsmtFinSF2'],
    axis=1, inplace=True)
```



```
In [538]: # Now check for other null/missing values
nullcols=housing_numeric.isnull().sum()
nullcols[nullcols>0]
```

```
Out[538]: MasVnrArea      869
BsmtFinSF1      467
BsmtUnfSF       118
TotalBsmtSF      37
2ndFlrSF       829
BsmtFullBath     856
FullBath         9
HalfBath        913
BedroomAbvGr      6
KitchenAbvGr      1
Fireplaces      690
GarageCars        81
GarageArea        81
WoodDeckSF       761
OpenPorchSF      656
EnclosedPorch    1252
dtype: int64
```

```
In [410]: #housing_numeric['EnclosedPorch'].median()
housing_numeric['BsmtFullBath'].mean()
```

```
Out[410]: 1.0281456953642385
```

```
In [539]: # impute the null/missing values with median of the corresponding data points
for col in nullcols.index:
    housing_numeric[col].fillna(housing_numeric[col].median(), inplace=True)
nullcols=housing_numeric.isnull().sum()
nullcols[nullcols>0]
```

```
Out[539]: Series([], dtype: int64)
```

There are few variables which has `year` column. We need to encode the year column to include that into the model creation, otherwise year values such as 1995, 2000, 1970 etc might effect the overall model performance. We can get a diff of these year column from the current year 2019 and process those

```
In [540]: # processing year columns
for yrCol in housing_numeric.columns:
    if str(yrCol).startswith('Year') or str(yrCol).startswith('Yr'):
        #print(yrCol)
        housing_numeric[yrCol]=housing_numeric[yrCol].astype(int)
        housing_numeric[yrCol+str('_diff_todate')]=housing_numeric[yrCol].a
        pply(lambda x: int(2019)-int(x))
        housing_numeric.drop(columns=[yrCol],axis=1,inplace=True)
housing_numeric.head()
```


Out[540]:

	MSSubClass	LotArea	OverallQual	OverallCond	MasVnrArea	BsmtFinSF1	BsmtUnfSF	TotalBsmtSF	1stFl
0	60	8450	7	5	196.0	706.0	150.0	856.0	
1	20	9600	6	8	203.0	978.0	284.0	1262.0	
2	60	11250	7	5	162.0	486.0	434.0	920.0	
3	70	9550	7	5	203.0	216.0	540.0	756.0	
4	60	14260	8	5	350.0	655.0	490.0	1145.0	

```
In [421]: # check the heatmap now for correlation of numeric variables
cor = housing_numeric.corr()
#cor
```

```
In [422]: # figure size
plt.figure(figsize=(20,15))

# heatmap
sns.heatmap(cor, cmap="YlGnBu", annot=True)
plt.show()
```



We can see SalesPrice has high correlation with variables like

- GarageArea: Size of garage in square feet
- GarageCars: Size of garage in car capacity
- GrLivArea: Above grade (ground) living area square feet
- 1stFlrSF: First Floor square feet
- 2ndFlrSF: Second floor square feet

There are multicollinearity as well like GarageCars highly correlated with GarageArea which is obvious

```
In [541]: # check the null/missing values in categorical variables
nullcols=housing_categorical.isnull().sum()
nullcols[nullcols>0]
```

```
Out[541]: Alley          1369
MasVnrType           8
BsmtQual            37
BsmtCond            37
BsmtExposure        38
BsmtFinType1        37
BsmtFinType2        38
Electrical           1
FireplaceQu        690
GarageType          81
GarageFinish        81
GarageQual          81
GarageCond          81
PoolQC             1453
Fence              1179
MiscFeature        1406
dtype: int64
```

```
In [542]: #drop the columns which has high count of null or missing values
housing_categorical.drop(columns=['Alley', 'Fence', 'PoolQC', 'MiscFeature'],
axis=1, inplace=True)
```

```
In [544]: nullcols=housing_categorical.isnull().sum()
nullcols[nullcols>0]
```

```
Out[544]: MasVnrType           8
BsmtQual            37
BsmtCond            37
BsmtExposure        38
BsmtFinType1        37
BsmtFinType2        38
Electrical           1
FireplaceQu        690
GarageType          81
GarageFinish        81
GarageQual          81
GarageCond          81
dtype: int64
```

Except Electrical all the other above variables can be imputed with no value like 'NA' or 'None'

```
In [545]: null_col_to_impute=nullcols[nullcols>1]
for col in null_col_to_impute.index:
    if col == "MasVnrType":
        housing_categorical[col].fillna("None", inplace=True)
    else:
        housing_categorical[col].fillna("NA", inplace=True)
```

```
In [546]: nullcols=housing_categorical.isnull().sum()
nullcols[nullcols>0]
```

```
Out[546]: Electrical          1
dtype: int64
```

```
In [547]: housing_categorical['Electrical'].value_counts()
```

```
Out[547]: SBrkr      1334  
FuseA        94  
FuseF        27  
FuseP         3  
Mix           1  
Name: Electrical, dtype: int64
```

```
In [548]: # Imputing the row with Unknown for missing value for Electrical variable  
#housing_categorical=housing_categorical[~np.isnan(housing_categorical['Electrical'])]  
housing_categorical['Electrical'].fillna('Unknown', inplace=True)
```

There are rating or quality related categorical variable whose values are having abbreviation like Ex for Excellent, Gd for Good etc. Where as the similar numerical variables like OverallQual, OverallCond has integer values for such ratings. We should encode such categorical variables like ExterQual, ExterCond, BsmtQual, BsmtCond etc in numeric format like OverallQual, OverallCond.

```
In [430]: housing_categorical['BsmtQual'].value_counts()
```

```
Out[430]: TA        649  
Gd         618  
Ex         121  
NA          37  
Fa          35  
Name: BsmtQual, dtype: int64
```

```
In [431]: housing_categorical['FireplaceQu'].value_counts()
```

```
Out[431]: NA        690  
Gd         380  
TA         313  
Fa          33  
Ex          24  
Po          20  
Name: FireplaceQu, dtype: int64
```

```
In [549]: print(housing_categorical.shape[1])  
print(housing_numeric.shape[1])
```

```
39  
28
```

```
In [550]: varlist = ['BsmtQual', 'BsmtCond', 'ExterQual', 'ExterCond',
                    'FireplaceQu', 'KitchenQual', 'HeatingQC']

# Defining the map function
def encode_cat_to_numeric(x):
    return x.map({'Ex': 9, "Gd": 7, "TA": 5, 'Fa': 3, 'NA': 0, 'Po': 2})

# Applying the function to the housing list
housing_numeric[varlist] = housing_categorical[varlist].apply(encode_cat_to_numeric)
for var in varlist:
    housing_numeric[var]=housing_numeric[var].astype(int)
    housing_categorical.drop([var], axis=1, inplace=True)

print(housing_categorical.shape[1])
print(housing_numeric.shape[1])
```

32
35

Data preparation

Now its time to process the categorical variables by creating dummies

```
In [552]: housing_dummies = pd.get_dummies(housing_categorical, drop_first=True)
housing_dummies.tail()
```

Out[552]:

	MSZoning_FV	MSZoning_RH	MSZoning_RL	MSZoning_RM	Street_Pave	LotShape_IR2	LotShape_IR3
1455	0	0	1	0	1	0	0
1456	0	0	1	0	1	0	0
1457	0	0	1	0	1	0	0
1458	0	0	1	0	1	0	0
1459	0	0	1	0	1	0	0

5 rows × 184 columns

```
In [553]: housing_dummies.shape
```

Out[553]: (1460, 184)

```
In [562]: # Concatenating all processed numerical and processed dummies for categoric
al
# This is the data set we will use going forward for model building
model_df = pd.concat([housing_numeric, housing_dummies], axis=1)
model_df.shape
```

Out[562]: (1460, 219)

```
In [561]: model_df.isnull().sum()
```

```
Out[561]: MSSubClass      14
          LotArea        14
          OverallQual    14
          OverallCond    14
          MasVnrArea     14
          BsmtFinSF1     14
          BsmtUnfSF      14
          TotalBsmtSF    14
          1stFlrSF       14
          2ndFlrSF       14
          GrLivArea      14
          BsmtFullBath   14
          FullBath       14
          HalfBath       14
          BedroomAbvGr   14
          KitchenAbvGr   14
          TotRmsAbvGrd   14
          Fireplaces     14
          GarageCars     14
          GarageArea     14
          WoodDeckSF     14
          OpenPorchSF    14
          EnclosedPorch  14
          MoSold         14
          SalePrice      14
          YearBuilt_diff_todate 14
          YearRemodAdd_diff_todate 14
          YrSold_diff_todate 14
          BsmtQual       14
          BsmtCond       14
          ExterQual       14
          ExterCond       14
          FireplaceQu     14
          KitchenQual     14
          HeatingQC       14
          MSZoning_FV     14
          MSZoning_RH     14
          MSZoning_RL     14
          MSZoning_RM     14
          Street_Pave     14
          LotShape_IR2    14
          LotShape_IR3    14
          LotShape_Reg     14
          LandContour_HLS  14
          LandContour_Low  14
          LandContour_Lvl  14
          Utilities_NoSeWa 14
          LotConfig_CulDSac 14
          LotConfig_FR2    14
          LotConfig_FR3    14
          LotConfig_Inside 14
          LandSlope_Mod    14
          LandSlope_Sev    14
          Neighborhood_Blueste 14
          Neighborhood_BrDale 14
          Neighborhood_BrkSide 14
          Neighborhood_ClearCr 14
          Neighborhood_CollgCr 14
          Neighborhood_Crawfor 14
          Neighborhood_Edwards 14
          Neighborhood_Gilbert 14
          Neighborhood_IDOTRR 14
          Neighborhood_MeadowV 14
```



```
In [564]: # because of the concat for some reason I'm getting 14 rows of NaN values i
n the merged data set model_df
# to adjust removing those rows
model_df=model_df[~np.isnan(model_df['MSSubClass'])]
#model_df.isnull().sum()
```

```
In [570]: model_df.shape
```

```
Out[570]: (1460, 219)
```

```
In [565]: # split X and y
y = model_df.loc[:, 'SalePrice']
X = model_df.loc[:, model_df.columns != 'SalePrice']
```

```
In [566]: # Scale the features using StnadardScaler
# scale
scaler = StandardScaler()
#scaler.fit(X)
scaler.fit_transform(X)
```

```
Out[566]: array([[ 0.07337496, -0.20714171,  0.65147924, ..., -0.11785113,
  0.4676514 , -0.30599503],
 [-0.87256276, -0.09188637, -0.07183611, ..., -0.11785113,
  0.4676514 , -0.30599503],
 [ 0.07337496,  0.07347998,  0.65147924, ..., -0.11785113,
  0.4676514 , -0.30599503],
 ...,
 [ 0.30985939, -0.14781027,  0.65147924, ..., -0.11785113,
  0.4676514 , -0.30599503],
 [-0.87256276, -0.08016039, -0.79515147, ..., -0.11785113,
  0.4676514 , -0.30599503],
 [-0.87256276, -0.05811155, -0.79515147, ..., -0.11785113,
  0.4676514 , -0.30599503]])
```

```
In [567]: # split into train and test set
X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    train_size=0.7,
                                                    test_size = 0.3, random
                                                    _state=100)
```

Model Building and Evaluation

Building models using both Ridge and Lasso technique

Ridge Regression

```
In [581]: def ridge_n_lasso(ml_algo):
# list of alphas to tune
    params = {'alpha': [0.0001, 0.001, 0.01, 0.05, 0.1,
0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0, 2.0, 3.0,
4.0, 5.0, 6.0, 7.0, 8.0, 9.0, 10.0, 20, 25, 40, 50]}

    # cross validation
    # folds = 5
    # creating a KFold object with 5 splits
    folds = KFold(n_splits = 5, shuffle = True, random_state = 100)

    model_ridge = GridSearchCV(estimator = ml_algo,
                                param_grid = params,
                                scoring= 'neg_mean_absolute_error',
                                cv = folds,
                                return_train_score=True,
                                verbose = 1)

    return model_ridge
```

```
In [582]: ridge = Ridge()
model_ridge=ridge_n_lasso(ridge)
model_ridge.fit(X_train, y_train)
```

Fitting 5 folds for each of 27 candidates, totalling 135 fits

[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.

[Parallel(n_jobs=1)]: Done 135 out of 135 | elapsed: 2.5s finished

```
Out[582]: GridSearchCV(cv=KFold(n_splits=5, random_state=100, shuffle=True),
    error_score='raise-deprecating',
    estimator=Ridge(alpha=1.0, copy_X=True, fit_intercept=True,
        max_iter=None, normalize=False, random_state=None,
        solver='auto', tol=0.001),
    iid='warn', n_jobs=None,
    param_grid={'alpha': [0.0001, 0.001, 0.01, 0.05, 0.1, 0.2, 0.3
,
        0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0, 2.0,
3.0,
        4.0, 5.0, 6.0, 7.0, 8.0, 9.0, 10.0, 20,
25,
        40, 50]},
    pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
    scoring='neg_mean_absolute_error', verbose=1)
```


```
In [583]: ridge_results = pd.DataFrame(model_ridge.cv_results_)
          ridge_results.head()
```

Out[583]:

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_alpha	params	split0_test_score	sp
0	0.030818	0.015069	0.006603	0.004493	0.0001	{'alpha': 0.0001}	-22456.054552	
1	0.032220	0.019469	0.003731	0.000899	0.001	{'alpha': 0.001}	-22453.360296	
2	0.012295	0.001765	0.003154	0.000602	0.01	{'alpha': 0.01}	-22427.176359	
3	0.008849	0.000643	0.002434	0.000126	0.05	{'alpha': 0.05}	-22295.478273	
4	0.009133	0.000930	0.002668	0.000248	0.1	{'alpha': 0.1}	-22114.346561	

```
In [584]: # plotting cv results
          #ridge_results['param_alpha'] = ridge_results['param_alpha'].astype('float')
          ridge_results['param_alpha'] = ridge_results['param_alpha'].astype('int32')

          # plotting
          plt.plot(ridge_results['param_alpha'], ridge_results['mean_train_score'])
          plt.plot(ridge_results['param_alpha'], ridge_results['mean_test_score'])
          plt.xlabel('alpha')
          plt.ylabel('Negative Mean Absolute Error')
          plt.title("Negative Mean Absolute Error and alpha")
          plt.legend(['train score', 'test score'], loc='upper right')
          plt.show()
```



We see from the plot that the test negative mean absolute error first increases and then follows a static trend. But the training negative mean absolute error keeps on decreasing as we increase the value of the hyperparameter, which is in accordance with the bias-variance trade-off.

Somewhere between $\alpha=8-10$ test performs peak, hence the final model we would generate with $\alpha=9$ we construct the objective function and re-run the regression algorithm on entire training data. Then we can have our final model with optimal features.

```
In [608]: alpha = 9
         ridge = Ridge(alpha=alpha)

         ridge.fit(X_train, y_train)
```

```
Out[608]: Ridge(alpha=9, copy_X=True, fit_intercept=True, max_iter=None, normalize=False,
              random_state=None, solver='auto', tol=0.001)
```

```
In [609]: # ridge predict
         y_train_pred = ridge.predict(X_train)
         print("ridge train pred r2_score: " + str(metrics.r2_score(y_true=y_train,
         y_pred=y_train_pred)))
         y_test_pred = ridge.predict(X_test)
         print("ridge test pred r2_score: " + str(metrics.r2_score(y_true=y_test, y_
         pred=y_test_pred)))
```

```
ridge train pred r2_score: 0.8826345775619088
ridge test pred r2_score: 0.866495729272287
```

```
In [610]: ridge_df=get_top_important_predictor_var(ridge)
         ridge_df.head(5)
```

```
Out[610]:
```

	feature	model_coefficients
67	Neighborhood_NridgHt	27644.272840
66	Neighborhood_NoRidge	25984.035574
113	RoofMatl_WdShngl	25917.171437
151	BsmtExposure_Gd	17864.974188
57	Neighborhood_Crawfor	14469.646771

```
In [611]: # doubling the alpha for Ridge
         alpha=18
         ridge=Ridge(alpha=alpha)
         ridge.fit(X_train,y_train)
```

```
Out[611]: Ridge(alpha=18, copy_X=True, fit_intercept=True, max_iter=None, normalize=False,
              random_state=None, solver='auto', tol=0.001)
```

```
In [612]: # ridge predict
         y_train_pred = ridge.predict(X_train)
         print("ridge train pred r2_score: " + str(metrics.r2_score(y_true=y_train,
         y_pred=y_train_pred)))
         y_test_pred = ridge.predict(X_test)
         print("ridge test pred r2_score: " + str(metrics.r2_score(y_true=y_test, y_
         pred=y_test_pred)))
```

```
ridge train pred r2_score: 0.8754148422544542
ridge test pred r2_score: 0.8658353012848261
```

```
In [613]: ridge_df=get_top_important_predictor_var(ridge)
         ridge_df.head(5)
```

Out[613]:

	feature	model_coefficients
67	Neighborhood_NridgHt	21550.333043
66	Neighborhood_NoRidge	20649.595561
151	BsmtExposure_Gd	16201.153533
113	RoofMatl_WdShngl	14563.197434
57	Neighborhood_Crawfor	11693.277195

```
In [575]: # ridge model parameters
         ridge_model_parameters=list(ridge.coef_)
         len(ridge_model_parameters)
         #ridge.intercept_
```

Out[575]: 218

Lasso Regression

```
In [614]: lasso=Lasso()
         model_lasso=ridge_n_lasso(lasso)
         model_lasso.fit(X_train, y_train)
```

Fitting 5 folds for each of 27 candidates, totalling 135 fits

[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.

[Parallel(n_jobs=1)]: Done 135 out of 135 | elapsed: 19.2s finished

```
Out[614]: GridSearchCV(cv=KFold(n_splits=5, random_state=100, shuffle=True),
                      error_score='raise-deprecating',
                      estimator=Lasso(alpha=1.0, copy_X=True, fit_intercept=True,
                                      max_iter=1000, normalize=False, positive=False),
                      ,
                      precompute=False, random_state=None,
                      selection='cyclic', tol=0.0001, warm_start=False),
          iid='warn', n_jobs=None,
          param_grid={'alpha': [0.0001, 0.001, 0.01, 0.05, 0.1, 0.2, 0.3,
                                0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0, 2.0,
                                3.0, 4.0, 5.0, 6.0, 7.0, 8.0, 9.0, 10.0, 20,
                                25, 40, 50]},
          pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
          scoring='neg_mean_absolute_error', verbose=1)
```


```
In [615]: lasso_results=pd.DataFrame(model_lasso.cv_results_)
lasso_results.head()
```

```
Out[615]:
```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_alpha	params	split0_test_score	sp
0	0.195279	0.136348	0.002692	0.000252	0.0001	{'alpha': 0.0001}	-22368.567357	
1	0.155852	0.041877	0.005612	0.005758	0.001	{'alpha': 0.001}	-22368.640773	
2	0.130968	0.004422	0.002580	0.000090	0.01	{'alpha': 0.01}	-22369.400146	
3	0.310048	0.140838	0.003560	0.001231	0.05	{'alpha': 0.05}	-22375.592085	
4	0.303399	0.140561	0.002726	0.000201	0.1	{'alpha': 0.1}	-22370.863167	

```
In [589]: lasso_results['param_alpha'] = lasso_results['param_alpha'].astype('int32')

# plotting
plt.plot(lasso_results['param_alpha'], lasso_results['mean_train_score'])
plt.plot(lasso_results['param_alpha'], lasso_results['mean_test_score'])
plt.xlabel('alpha')
plt.ylabel('Negative Mean Absolute Error')
plt.title("Negative Mean Absolute Error and alpha")
plt.legend(['train score', 'test score'], loc='upper right')
plt.show()
```



```
In [616]: # select the optimal alpha as 9
alpha = 9
lasso = Lasso(alpha=alpha)
lasso.fit(X_train, y_train)
```

```
Out[616]: Lasso(alpha=9, copy_X=True, fit_intercept=True, max_iter=1000, normalize=False,
              positive=False, precompute=False, random_state=None, selection='cyclic',
              tol=0.0001, warm_start=False)
```

```
In [617]: # lasso predict
y_train_pred = lasso.predict(X_train)
print("lasso train pred r2_score: " + str(metrics.r2_score(y_true=y_train,
y_pred=y_train_pred)))
y_test_pred = lasso.predict(X_test)
print("lasso test pred r2_score: " + str(metrics.r2_score(y_true=y_test, y_
pred=y_test_pred)))
```

```
lasso train pred r2_score: 0.9343831436100587
lasso test pred r2_score: 0.8312355842656285
```

There is a significant difference between training and test data set `r2_score` in Lasso regression. With `lambda=9` (`alpha=9`) Ridge model suggests `r2_score` is 86% which is much closer match the training data set `r2_score`. But Lasso model suggests `r2_score` is only 83% where as training data set `r2_score` is 93%. This shows Lasso model is underfit with `alpha=9`

```
In [618]: # top 5 predictor variable for lasso when alpha=9
lasso_df=get_top_important_predictor_var(lasso)
lasso_df.head(5)
```

Out[618]:

	feature	model_coefficients
113	RoofMatl_WdShngl	613913.194110
108	RoofMatl_Membran	566075.792513
107	RoofMatl_CompShg	543666.861094
109	RoofMatl_Metal	539081.567614
111	RoofMatl_Tar&Grv	532691.993097

```
In [619]: # adjusting the alpha to make it double
alpha = 18
lasso = Lasso(alpha=alpha)
lasso.fit(X_train, y_train)

y_train_pred = lasso.predict(X_train)
print("lasso train pred r2_score: " + str(metrics.r2_score(y_true=y_train,
y_pred=y_train_pred)))
y_test_pred = lasso.predict(X_test)
print("lasso test pred r2_score: " + str(metrics.r2_score(y_true=y_test, y_
pred=y_test_pred)))
```

```
lasso train pred r2_score: 0.9279924910225573
lasso test pred r2_score: 0.8393765560695181
```

```
In [620]: # top 5 predictor variable for lasso when alpha=18
lasso_df=get_top_important_predictor_var(lasso)
lasso_df.head(5)
```

Out[620]:

	feature	model_coefficients
113	RoofMatl_WdShngl	464322.606528
107	RoofMatl_CompShg	395235.659841
108	RoofMatl_Membran	376521.872650
111	RoofMatl_Tar&Grv	370027.574375
110	RoofMatl_Roll	368979.404766

```
In [603]: # adjusting the alpha to make it 18*2=36
alpha = 36
lasso = Lasso(alpha=alpha)
lasso.fit(X_train, y_train)

y_train_pred = lasso.predict(X_train)
print("lasso train pred r2_score: " + str(metrics.r2_score(y_true=y_train,
y_pred=y_train_pred)))
y_test_pred = lasso.predict(X_test)
print("lasso test pred r2_score: " + str(metrics.r2_score(y_true=y_test, y_
pred=y_test_pred)))

lasso train pred r2_score: 0.9079437886596565
lasso test pred r2_score: 0.8454155310419142
```

```
In [604]: # top 5 predictor variable for lasso when alpha=18*2=36
lasso_df=get_top_important_predictor_var(lasso)
lasso_df.head(5)
```

Out[604]:

	feature	model_coefficients
113	RoofMatl_WdShngl	186433.733644
107	RoofMatl_CompShg	117564.346121
112	RoofMatl_WdShake	81407.903798
111	RoofMatl_Tar&Grv	75942.509302
110	RoofMatl_Roll	68375.171794

```
In [621]: # adjusting the alpha 36*2=72
alpha = 72
lasso = Lasso(alpha=alpha)

lasso.fit(X_train, y_train)

y_train_pred = lasso.predict(X_train)
print("lasso train pred r2_score: " + str(metrics.r2_score(y_true=y_train,
y_pred=y_train_pred)))
y_test_pred = lasso.predict(X_test)
print("lasso test pred r2_score: " + str(metrics.r2_score(y_true=y_test, y_
pred=y_test_pred)))

lasso train pred r2_score: 0.8959380007832307
lasso test pred r2_score: 0.8563338291209158
```


With alpha=72 model test r2_score 85% close matches (< than 5 degree diff) the training r2_score which is close to 90%. But the training r2_score fall below 90% beyond alpha > 40

```
In [622]: # top 5 predictor variable for lasso when alpha=36*2=72
lasso_df=get_top_important_predictor_var(lasso)
lasso_df.head(5)
```

Out[622]:

	feature	model_coefficients
113	RoofMatl_WdShngl	90101.229661
67	Neighborhood_NridgHt	38380.891299
66	Neighborhood_NoRidge	37301.698127
107	RoofMatl_CompShg	30448.521618
73	Neighborhood_StoneBr	23902.282064

```
In [260]: # Check the lasso model parameters
lasso_model_parameters = list(lasso.coef_)
lasso_model_parameters.insert(0, lasso.intercept_)
lasso_model_parameters = [round(x, 3) for x in lasso_model_parameters]
cols = X.columns
cols = cols.insert(0, "constant")
list(zip(cols, lasso_model_parameters))
```

```
Out[260]: [('constant', -93901.212),
('MSSubClass', -200.91),
('LotArea', 0.335),
('OverallQual', 12148.541),
('OverallCond', 4500.622),
('YearBuilt', 194.274),
('YearRemodAdd', 104.091),
('MasVnrArea', 21.75),
('BsmtFinSF1', -1.435),
('BsmtUnfSF', -10.757),
('TotalBsmtSF', 12.546),
('1stFlrSF', -6.92),
('2ndFlrSF', 25.388),
('GrLivArea', 60.894),
('BsmtFullBath', -3221.848),
('FullBath', 1736.238),
('HalfBath', -18005.305),
('BedroomAbvGr', -3723.222),
('KitchenAbvGr', -11719.015),
('TotRmsAbvGrd', 1554.766),
('Fireplaces', 3910.191),
('GarageCars', 13635.243),
('GarageArea', -10.996),
('WoodDeckSF', 15.64),
('OpenPorchSF', 22.918),
('EnclosedPorch', 6.323),
('MoSold', -273.451),
('YrSold', -313.249),
('MSZoning_FV', 0.0),
('MSZoning_RH', 2103.465),
('MSZoning_RL', 6961.223),
('MSZoning_RM', 1012.098),
('Street_Pave', 9713.479),
('LotShape_IR2', 0.0),
('LotShape_IR3', -29884.569),
('LotShape_Reg', 2369.94),
('LandContour_HLS', 13142.982),
('LandContour_Low', 5847.344),
('LandContour_Lvl', 9520.145),
('Utilities_NoSeWa', -0.0),
('LotConfig_CulDSac', 14187.242),
('LotConfig_FR2', -7093.561),
('LotConfig_FR3', -0.0),
('LotConfig_Inside', 276.884),
('LandSlope_Mod', 570.881),
('LandSlope_Sev', 10443.869),
('Neighborhood_Blueste', 0.0),
('Neighborhood_BrDale', 7654.184),
('Neighborhood_BrkSide', 4656.396),
('Neighborhood_ClearCr', 2753.336),
('Neighborhood_CollgCr', -992.442),
('Neighborhood_Crawfor', 23356.547),
('Neighborhood_Edwards', -9939.022),
('Neighborhood_Gilbert', -2051.484),
('Neighborhood_IDOTRR', -1533.103),
('Neighborhood_MeadowV', 0.0),
('Neighborhood_Mitchel', -9320.948),
('Neighborhood_NAmes', -5684.286),
('Neighborhood_NPkVill', 0.0),
('Neighborhood_NWAmes', -6586.74),
('Neighborhood_NoRidge', 35100.793),
('Neighborhood_NridgHt', 44521.819),
('Neighborhood_OldTown', -2328.441),
```

```
In [607]: def get_top_important_predictor_var(regModel):  
    cols = X.columns  
    #cols  
    model_parameters = list(regModel.coef_)  
    df=pd.DataFrame(list(zip(cols, model_parameters)))  
  
    df['feature']=df[0]  
    df['model_coefficients']=df[1]  
    df.drop([0,1],inplace=True,axis=1)  
    df=df.sort_values(by='model_coefficients',ascending=False)  
    return df
```